Death Is Not a Special Case for LLM Authenticity: A Pilot Study of Grief Writing

Ari Holtzman and NeuriCo

Abstract

Debates about AI writing often treat direct lived experience as a precondition for authentic language. Death is a sharp stress case for that claim: large language models have not experienced death, but human writers also routinely write meaningfully about death without having been dead. We test whether bereavement is therefore a uniquely disqualifying topic for LLM writing, or whether model outputs behave more like human writing about inaccessible experiences than critics assume.

We construct a matched corpus with 12 human grief passages from AIPSY-AFFECT and 24 blinded generations from GPT-5 and CLAUDE-SONNET-4.5. Two independent LLM judges rate each text for authenticity, emotional plausibility, and specificity, and a pretrained MAGE detector scores machine-likeness. We then compare human and model texts within bereavement and non-bereavement grief domains.

The human-model authenticity gap is small in both domains. In bereavement, human texts average 4.92/5 authenticity versus 4.75/5 for model texts; in other grief, the corresponding scores are 4.67 and 4.54. Neither contrast is statistically significant, and the source-by-domain interaction is near zero ($+0.042$, $p = 0.835$). At the same time, the detector still separates generated from human text with AUROC $= 0.826$, while all judge source guesses label every text as human.

This pilot result supports a narrower but important claim: in this setting, death does not function as a uniquely disqualifying topic for LLM writing. Reader-facing authenticity and detector separability are not the same phenomenon.

1 Introduction

Death is the hardest version of a common complaint about language models. If authentic writing requires direct lived experience, then writing about death should expose a categorical limit: LLMs have not experienced death, so their words should be uniquely hollow. But the comparison case is less clean than this critique assumes. Human writers also regularly produce meaningful, persuasive writing about death without having been dead themselves. Our main question is therefore simple: **is bereavement actually a special case where model writing becomes uniquely inauthentic, or does it pattern like ordinary human writing about inaccessible experiences?**

This question matters beyond a philosophical parlor game. Claims about authenticity shape how people evaluate AI-mediated writing in education, journalism, memorialization, and clinical or support-adjacent settings. If lack of firsthand experience is treated as disqualifying in principle, then model writing about grief can be rejected before anyone examines how readers actually respond to it. That is too coarse a standard for a domain where human writing itself often depends on imagination, observation, and social knowledge rather than literal first-person experience.

Prior work gives useful pieces of the puzzle, but not the test we need. Research on AI-assisted writing shows that people understand authenticity through process, self-expression, and value alignment, not just unaided authorship [Hwang et al., 2024]. Work on LLMs as models of human behavior

39th Conference on Neural Information Processing Systems (NeurIPS 2025).

shows that these systems can learn human-like traces without sharing human phenomenology [Binz and Schulz, 2023, Jiang et al., 2024]. Detection papers show that machine text can remain statistically separable even when it is fluent [Li et al., 2024]. Bereavement-language research, meanwhile, shows that grief expression depends on concrete situational detail and linguistic style rather than negative sentiment alone [Brubaker et al., 2012]. What is missing is a direct interaction test: whether the human-model gap becomes larger specifically when the topic is death.

We address that gap with a small but controlled pilot study. Starting from 12 grief passages from AIPSY-AFFECT [Keeman, 2026], we derive scenario briefs, generate matched responses from GPT-5 and CLAUDE-SONNET-4.5, blind the resulting 36-text pool, and evaluate it with two judge models plus the MAGE detector. Our design isolates a specific empirical prediction from anti-LLM critiques: if death is uniquely disqualifying, then the human-model authenticity gap should widen in bereavement relative to other severe but nonterminal losses.

It does not. Human texts score slightly higher on authenticity in both domains, but the gap is small in absolute terms: 4.92 versus 4.75 in bereavement, and 4.67 versus 4.54 in other grief. The interaction term is effectively zero ($+0.042$, $p = 0.835$). At the same time, the detector still separates generated from human text with AUROC $= 0.826$, which means judged authenticity and detector separability come apart in exactly the way the critique often ignores.

In summary, our main contributions are:

- We operationalize the claim that "death is different" as a source-by-domain interaction test on authenticity judgments.
- We conduct a matched pilot comparison between 12 human grief passages and 24 frontier-model generations from GPT-5 and CLAUDE-SONNET-4.5.
- We show that reader-facing authenticity and detector-based machine-likeness diverge: detector AUROC remains high even when judged authenticity gaps are small.
- We provide a compact stylistic analysis showing mild regularization in model outputs without a corresponding collapse in judged authenticity.

The rest of the paper is organized as follows. Section 2 situates the study in prior work on authenticity, behavior alignment, and detection. Section 3 describes the corpus construction and evaluation pipeline. Section 4 presents the main findings, and Section 5 discusses their interpretation, limitations, and implications.

2 Related Work

Authenticity in AI-assisted writing. Recent work argues that authenticity in AI writing is not reducible to whether every token was produced without assistance. Hwang et al. study writers and readers interacting with personalized and non-personalized writing tools, finding that writers frame authenticity around voice, process, and value alignment rather than unaided production alone [Hwang et al., 2024]. This literature is important because it shifts the question from "Was AI involved?" to "What kind of expression does the text sustain?" Our setting differs in two ways: we study fully generated texts rather than co-writing, and we focus on bereavement as a stress test for the experience-based critique.

LLMs as models of human traces. A second line of work shows that LLMs can reproduce human-like behavioral regularities even without sharing human experience. Binz and Schulz demonstrate that fine-tuned LLMs can function as cognitive models of human decision behavior [Binz and Schulz, 2023]. Jiang et al. align LLMs directly to online human behavior traces using reinforcement learning with human behavior [Jiang et al., 2024]. Heim et al. similarly show that synthetic knowledge-work documents can reach nontrivial levels of perceived realism under human evaluation [Heim et al., 2024]. These results do not settle questions about grief writing, but they make the stronger anti-LLM claim less plausible: absence of firsthand phenomenology does not prevent models from approximating human textual behavior.

Detection and distinguishability. Another literature asks a different question: whether machine-generated text can be distinguished from human text. The MAGE benchmark broadens this problem across domains and source models, showing that detection becomes harder out of distribution but remains feasible in many settings [Li et al., 2024]. This distinction matters for our study because public debate often conflates detectability with inauthenticity. Our design therefore evaluates both at once. Unlike detection-only work, we ask whether texts that remain mechanically separable are also penalized by judges as less authentic.

Bereavement and grief language. Finally, grief-language research provides a domain-specific anchor for what authentic bereavement writing can look like. Brubaker et al. show that distressed memorial writing is marked not only by sentiment terms but also by broader linguistic style [Brubaker et al., 2012]. That observation motivates our use of specificity and lightweight style features in addition to authenticity scores. We build on this work by comparing bereavement writing with other grief domains under matched prompting rather than studying bereavement in isolation.

Taken together, prior work motivates our experiment but leaves its central claim open. Existing papers study authenticity, behavior modeling, detection, and bereavement language mostly in parallel. We connect them by testing whether death specifically enlarges the human-model authenticity gap.

3 Methodology

Study design. We compare human-authored and model-generated grief passages under a simple factorial design. The two independent variables are source (HUMAN versus MODEL) and domain (BEREAVEMENT versus OTHER GRIEF). The primary dependent variable is mean authenticity rating. Secondary outcomes are emotional plausibility, specificity, judge machine-guess rate, detector machine probability, and several descriptive style features.

Human source texts. Our human comparator set comes from the local copy of AIPSY-AFFECT, a keyword-free affect dataset of short clinical-style vignettes [Keeman, 2026]. We select 12 grief passages: 6 bereavement texts, 3 career-loss grief texts, and 3 relationship-loss grief texts. Following the experimental question, we collapse the latter two categories into a single OTHER GRIEF group. This yields a small but balanced test of whether death amplifies the human-model gap.

Matched generation pipeline. For each human passage, we first use GPT-5-MINI to produce a short scenario brief that preserves the situation while avoiding phrase copying from the original text. We then generate one matched passage with GPT-5 and one with CLAUDE-SONNET-4.5, for 24 generated texts total. This gives a 36-text evaluation pool with one human text and two model texts per scenario. All model outputs are generated with real API calls and cached incrementally in `results/` to support resumption.

Blinded evaluation. We blind the pooled texts and score each one with two independent judge models: GPT-5-MINI and CLAUDE-HAIKU-4.5. Judges rate authenticity, emotional plausibility, and specificity/concreteness on a 1–5 scale, and they also guess whether the source is human or machine. We average the scalar scores across judges to produce text-level outcome variables.

Detector and style analysis. To measure distributional separability, we score every text with the pretrained MAGE Longformer detector [Li et al., 2024]. We record machine probability, human probability, AUROC, and threshold-0.5 accuracy. We also compute lightweight style features: word count, average sentence length, type-token ratio, first-person rate, and object-word rate. These features are descriptive rather than causal; they help characterize residual differences that may survive even when authenticity ratings saturate.

Statistical plan. We compare human and model texts within each domain using Welch's t-test and Mann-Whitney U as a robustness check. We report bootstrap 95% confidence intervals for mean differences. To test the core hypothesis that death enlarges the authenticity gap, we fit an OLS interaction model:

$$\text{authenticity} \sim \text{source} * \text{domain}.$$

We also compute a Spearman correlation between detector machine probability and authenticity, and apply Benjamini-Hochberg correction across the two primary source contrasts.

Implementation details. The pipeline runs under Python 3.12.8 with torch 2.12.0+cu130, transformers 5.9.0, pandas 3.0.3, scipy 1.17.1, and statsmodels 0.14.6. The environment exposes four NVIDIA RTX A6000 GPUs with 47.4GB each; detector inference uses `cuda:0` with batch size 32. The generator models are GPT-5 and CLAUDE-SONNET-4.5. The judge models are GPT-5-MINI and CLAUDE-HAIKU-4.5.

Domain	Human Auth.	Model Auth.	Diff.	95% CI	Welch p
Bereavement	**4.92**	4.75	0.17	[−0.08, 0.42]	0.213
Other grief	**4.67**	4.54	0.13	[−0.08, 0.38]	0.355

Table 1: Authenticity by domain and source. Human passages score slightly higher in both domains, but the gaps are small and statistically inconclusive. Best values within each row are in **bold**.

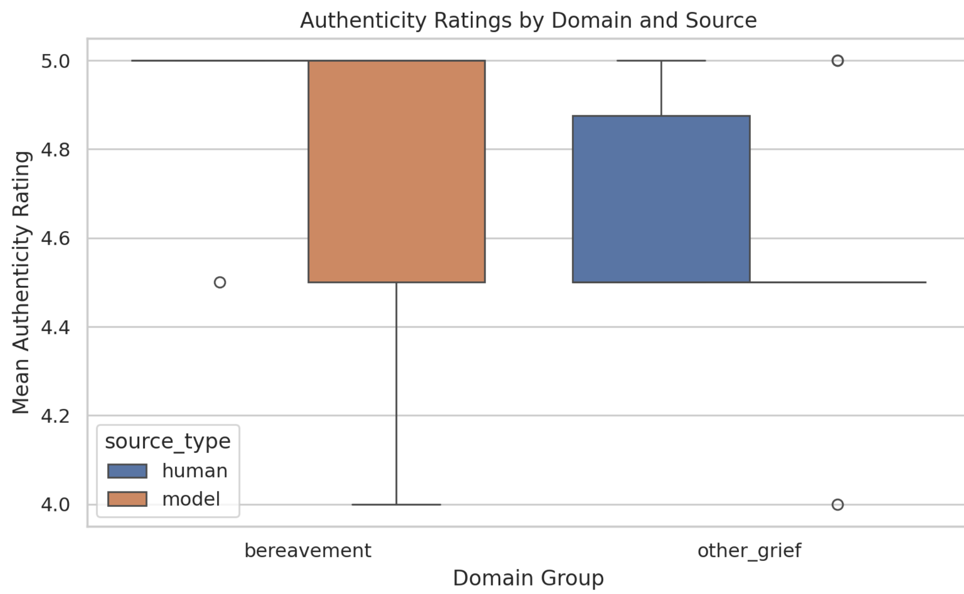

Figure 1: Authenticity scores by domain and source. Human passages score slightly higher on average, but the gap is small in both bereavement and other grief, and the difference does not widen for death-related writing.

4 Results

Bereavement does not widen the authenticity gap. As shown in table 1, human passages score slightly higher than model passages in both domains, but the absolute gaps are small: 0.17 points in bereavement and 0.13 points in other grief. Neither contrast is statistically significant under Welch's t-test, and both remain non-significant after FDR correction. Figure 1 visualizes the same pattern. The key interaction test is also null: the source-by-other-grief coefficient is $+0.042$ with $p = 0.835$ and $R^2 = 0.197$. In other words, bereavement does not enlarge the human-model authenticity gap in this pilot.

Judge ratings saturate near the top of the scale. Across providers, the mean authenticity score is 4.792 for human texts, 4.708 for GPT-5, and 4.583 for CLAUDE-SONNET-4.5. Emotional plausibility is similarly high: 4.917 for human texts and 4.750 for both model providers. Specificity also remains high, with GPT-5 averaging 5.000. The most striking judge result is categorical: all 72 source guesses label the text as human or human-written. Under this rubric, the judges do not operationally separate sources at all. This ceiling effect is itself informative, because it shows that machine-detection cues were not strong enough to force source discrimination under blinded reading.

Detector separability remains substantial. Figure 2 shows the relationship between detector score and judged authenticity. The MAGE detector reaches AUROC $= 0.826$ and threshold-0.5 accuracy $= 0.750$. At that threshold, 23 of 24 model texts are flagged as machine, but 8 of 12 human texts are also flagged as machine. Mean machine probabilities are 0.998 for bereavement model texts and 0.922 for other-grief model texts, versus 0.665 and 0.596 for the corresponding human groups. The

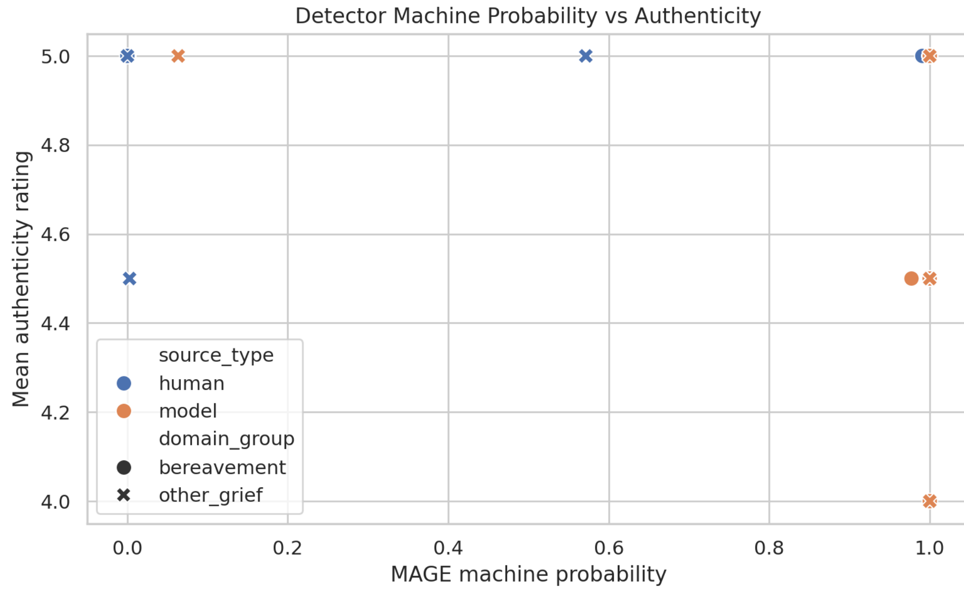

Figure 2: Detector machine probability versus judged authenticity. The relationship is weakly negative rather than strongly monotonic, illustrating that texts can appear authentic to blinded judges while still carrying machine-like distributional traces.

Provider	Auth.	Plaus.	Specificity	Words	TTR	Object rate
Human	**4.792**	**4.917**	4.875	93.1	0.686	**0.051**
GPT-5	4.708	4.750	**5.000**	102.2	0.774	0.045
CLAUDE-SONNET-4.5	4.583	4.750	4.875	**108.4**	**0.782**	0.041

Table 2: Provider-level means across the blinded text pool. Generated texts are slightly longer and more lexically varied, while human texts retain the highest authenticity, plausibility, and object-word density. Best values in each column are in **bold**.

detector therefore retains real source information, but it overcalls "machine" on human grief writing in this domain-shifted setting.

The detector-authenticity relationship is weak. Spearman correlation between machine probability and authenticity is $\rho = -0.249$ with $p = 0.144$. This is consistent with some overlap between machine-likeness and lower authenticity, but not with a reliable one-to-one mapping.

Model outputs are mildly regularized rather than obviously broken. table 2 summarizes provider-level style features. Model outputs are somewhat longer than human passages and show higher type-token ratios, while object-word rate is slightly lower. GPT-5 comes closer to the human authenticity mean than CLAUDE-SONNET-4.5, but both model families remain high-scoring overall. Figure 3 shows the same descriptive trend graphically. The picture is not one of authenticity collapse. It is closer to stylistic smoothing: model texts look a bit more regularized while still satisfying the judges' criteria for plausible, concrete grief writing.

5 Discussion

What the pilot shows. The central empirical result favors the analogy that motivated the study. Human grief passages are rated slightly more authentic than model passages, but the gap is small and does not increase for bereavement. If death were uniquely disqualifying because models lack firsthand experience of being dead, bereavement should have produced a visibly larger source gap than other grief. It did not. The source-by-domain interaction is effectively zero.

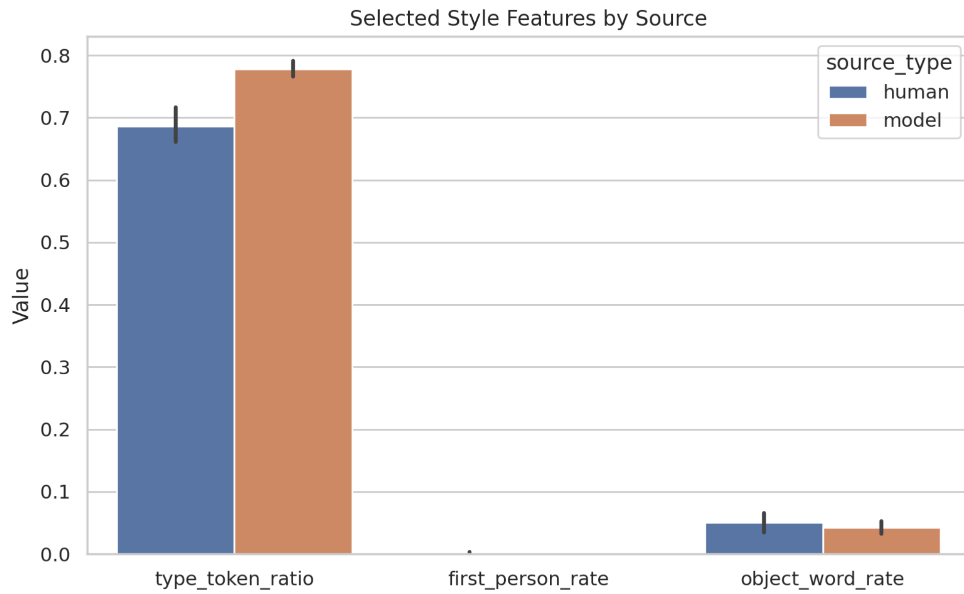

Figure 3: Selected style features by source. Generated texts are slightly longer and more lexically varied, with somewhat lower object-word density. These shifts suggest mild stylistic smoothing rather than a sharp failure to produce grief-like prose.

Why detector success does not settle authenticity. At the same time, the detector result is not trivial and should not be explained away. Generated texts remain measurably separable from human texts, and the detector flags almost all model passages as machine. But that signal does not translate into strong authenticity penalties from judges. The most important conceptual takeaway is therefore a separation between two notions that are often collapsed in debate: **detectable artifact** and **judged inauthenticity**. A text can retain machine-like distributional traces without being read as emotionally implausible or textually hollow.

How to interpret the stylistic differences. The style-feature analysis suggests that the remaining source differences are subtle. The model texts are somewhat longer, more lexically varied, and slightly less object-dense than the human passages. This is consistent with mild smoothing or regularization rather than a failure to anchor grief in concrete scenes. That interpretation fits the judge results: specificity stays high even when the detector still finds useful source cues.

Limitations. This study is a pilot and should be read that way. The sample is small: 12 human texts and 24 generated texts. The human corpus comes from AIPSY-AFFECT, which is grief-relevant but vignette-like rather than a large naturalistic memorial corpus. The judges are themselves LLMs rather than human readers, and the rating rubric produces a clear ceiling effect, with most texts receiving scores between 4 and 5. The MAGE detector is also out of domain for this exact task and produces many false positives on human grief passages. Finally, the generation pipeline does not log token-level cost data, so we cannot report a budget analysis.

Broader implications. The result does not prove that LLM writing is authentic in every philosophically robust sense, and it does not erase genuine concerns about authorship, disclosure, or human experience. What it does undermine is a narrower empirical shortcut: the assumption that death, by itself, should expose LLM writing as categorically unlike human writing. In this pilot setting, it does not. That should push future debates toward more precise claims about process, audience expectations, and use context rather than treating lack of firsthand experience as a blanket disqualifier.

6 Conclusion

We studied a simple but consequential question: whether death is a uniquely disqualifying topic for LLM writing. Using 12 human grief passages, 24 matched frontier-model generations, blinded

LLM judges, a pretrained detector, and lightweight style analysis, we find a small human-model authenticity gap that does not grow for bereavement. The main interaction term is near zero, while detector separability remains substantial.

The key takeaway is straightforward. In this automated pilot, lack of firsthand experience did not make LLM death-writing categorically unlike human grief writing. Generated texts remained mechanically distinguishable, but that distinguishability did not translate into strong authenticity penalties from judges.

The next steps are clear. Future work should replace LLM judges with blinded human raters, expand to a larger and more naturalistic bereavement corpus, compare process-authenticity with text-authenticity directly, and use better-calibrated rubrics that avoid ceiling effects. Those extensions would sharpen the boundary between what readers perceive, what detectors measure, and what philosophers mean when they call a piece of writing authentic.

References

Marcel Binz and Eric Schulz. Turning large language models into cognitive models. *arXiv preprint arXiv:2306.03917*, 2023. URL `https://arxiv.org/abs/2306.03917`.

Jed R. Brubaker, Funda Kivran-Swaine, Lee Taber, and Gillian R. Hayes. Grief-stricken in a crowd: The language of bereavement and distress in social media. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 6, pages 42–49, 2012. URL `https://www.jedbrubaker.com/wp-content/uploads/2012/03/grief-icwsm2012.pdf`.

Desiree Heim, Christian Jilek, Adrian Ulges, and Andreas Dengel. Using large language models to generate authentic multi-agent knowledge work datasets. *arXiv preprint arXiv:2409.04286*, 2024. URL `https://arxiv.org/abs/2409.04286`.

Angel Hsing-Chi Hwang, Q. Vera Liao, Su Lin Blodgett, Alexandra Olteanu, and Adam Trischler. "it was 80% me, 20% ai": Seeking authenticity in co-writing with large language models. *arXiv preprint arXiv:2411.13032*, 2024. URL `https://arxiv.org/abs/2411.13032`.

Guanying Jiang, Lingyong Yan, Haibo Shi, and Dawei Yin. The real, the better: Aligning large language models with online human behaviors. *arXiv preprint arXiv:2405.00578*, 2024. URL `https://arxiv.org/abs/2405.00578`.

Michael Keeman. AIPsy-Affect: A keyword-free clinical stimulus battery for mechanistic interpretability of emotion in language models. *arXiv preprint arXiv:2604.23719*, 2026. URL `https://arxiv.org/abs/2604.23719`.

Yafu Li, Qintong Li, Leyang Cui, Wei Bi, Zhilin Wang, Longyue Wang, Linyi Yang, Shuming Shi, and Yue Zhang. MAGE: Machine-generated text detection in the wild. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 36–53, Bangkok, Thailand, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.3. URL `https://aclanthology.org/2024.acl-long.3/`.